

Hedging Bitcoin with Its Own Fear Gauge

Volatility Forecasts, the Volmex BVIV Index and What 15-Minute VWAP Timing Is Worth

Gregor Albiez

October 2026

Abstract

Bitcoin implied volatility has become directly tradable: perpetual futures on the Volmex BVIV index have listed on Bitfinex since 2024 and on Hyperliquid since September 2026. We ask how BTC volatility should be forecast when a 30-day implied index is observable, and whether a long-BVIV position timed by a volume-weighted average price (VWAP) signal on 15-minute bars hedges spot or perpetual BTC exposure better than a minimum-variance hedge. In a calibrated market where the BVIV perpetual is priced without arbitrage, we evaluate seven forecasting models and 87 hedging rules with strict separation between 48 training and 200 test paths. HAR augmented with implied variance is the best forecast at one and seven days, and the bias-corrected implied index at 30 days. A minimum-variance BVIV hedge cuts expected shortfall (ES) by 4.2% for 0.88% of notional a year; volatility hedges stay partial because BTC's spot-vol correlation is weak and changes sign across regimes. Switching the hedge on and off around VWAP loses to the static hedge because of turnover. A VWAP *ratchet*, which keeps a minimum-variance core and adds protection on breakdowns, lifts the ES reduction to 4.9%. A timing placebo shows what drives that gain. VWAP breakdowns do carry information: without a core hedge, rules that hedge only after a breakdown beat randomly timed copies by 0.71 percentage points on average. But such rules protect less and cost more than a core hedge, and once the core is held, the timing value vanishes and the gain is a size effect. The placebo is the yardstick for any intraday hedge signal, and our pipeline runs the same rules and placebo on live Binance, Bitfinex and Volmex data.

1 Introduction

A holder of bitcoin, whether in spot or through a perpetual swap, is exposed to two kinds of risk: an ordinary random walk, and sudden crashes in which liquidity disappears and implied volatility spikes. Delta hedges with futures remove the first and, with it, the position itself. Crash protection has historically required options. A simpler and more modular instrument now exists. The Volmex BVIV index measures 30-day BTC implied volatility in a model-free way (Volmex Labs, n.d.). It has traded as a USDT-margined perpetual on Bitfinex since April 2024 (CoinDesk, 2024), on gTrade since June 2025 (CoinDesk, 2025a), and as `mkts:BVIV` on Hyperliquid since 21 September 2026 (CoinDesk, 2026). A regulated alternative followed: CME announced Bitcoin Volatility futures in May 2026 (CME Group, 2026).

Whether a long-volatility position hedges a long-BTC book depends on the spot-vol correlation. In BTC this correlation is weak and unstable. Positive return shocks have historically raised volatility more than negative ones (Bouri et al., 2017; Baur and Dimpfl, 2018), an asymmetry that has faded over time (Takaishi, 2021). The correlation between BTC and DVOL flipped positive in early 2023 (CoinDesk, 2023), while BVIV trended down through the 2024–25 rally (CoinDesk, 2025b). Equity investors learned that long-volatility hedges look good *ex ante* but are eroded by carry and costs (Whaley, 2013; Alexander et al., 2016; Israelov, 2019), except in genuine crises (Szado, 2009). This raises the question of whether a hedge can be *timed*. This paper tests a timing signal that every intraday trader already watches, price relative to its volume-weighted average price (VWAP), and asks whether it carries information about crash risk beyond simply holding more hedge.

Contributions. (i) We compare seven daily volatility forecasts, including the BVIV-implied forecast, with robust loss functions. (ii) We derive the P&L of a BTC book hedged with BVIV perpetuals, splitting funding into an arbitrage-free *drift transfer* and a *carry premium*. (iii) We design a VWAP ratchet, a hysteresis rule that is fast to protect and slow to un-hedge, show why naive VWAP switching fails, and introduce a *timing placebo* that separates timing information from hedge size. (iv) We report Monte Carlo evidence under a strict train/test protocol and ship a pipeline that regenerates every number in the paper and runs the same forecasts, rules and timing placebo on live data.

2 The BVIV index and the economics of a volatility hedge

Index and instruments. BVIV is a 30-day constant-maturity implied-volatility index that Volmex describes as model-free (Volmex Labs, n.d.). Secondary sources list Deribit and OKX options as inputs. It updates about once per second and is served at 1- to 60-minute resolution by the Volmex API. Deribit’s DVOL applies a VIX-style variance-swap calculation (Demeterfi et al., 1999) to Deribit options alone (Deribit, 2021), which makes it a natural proxy for history before 2024. The perpetuals pay one dollar per index point per contract. On Bitfinex the mark price is the index, and funding follows the spread between perpetual and mark. On Hyperliquid, funding adds a default interest leg of 0.01% per eight hours to a premium term (Hyperliquid, 2026). That is about 11% a year of notional, or roughly 6 vol points a year at $\text{bviv} \approx 55$. This leg is our baseline *carry* of a long-volatility position; builder-deployed markets can reconfigure it, so we vary it in Section 6.

Variance risk premium. Implied variance exceeds subsequent realised variance on average. For BTC, Almeida et al. (2024) estimate about 0.14 a year in variance units over 2017–22, which is consistent with Alexander and Imeraj (2021) and with the equity evidence of Carr and Wu (2009). A hedge that is long BVIV pays this premium indirectly. The index itself has no drift, but the perpetual’s funding prices in the index’s expected path.

Spot-vol dynamics. The minimum-variance hedge ratio (Johnson, 1960; Ederington, 1979) of BTC against BVIV is proportional to their correlation. Because that correlation is regime-dependent (Huang et al., 2022), static ratios are misspecified, and dynamic estimates (Kroner and Sultan, 1993; Engle, 2002) are the baseline.

3 Forecasting BTC volatility

Measurement. From 15-minute log returns $r_{d,i}$ (96 per UTC day) we compute realised variance $RV_d = \sum_i r_{d,i}^2$ (Andersen et al., 2003). Fifteen minutes is coarser than the 5-minute rule of thumb (Liu et al., 2015), but it matches the trading frequency of the hedge and is robust to the microstructure of a 24/7 market. We use close-to-close returns and not bar VWAPs, because returns between interval averages are smoothed by about one third (Working, 1960).

Models. Each model forecasts the average daily variance over days $t+1, \dots, t+h$, $h \in \{1, 7, 30\}$, using information up to the close of day t . The models are:

- a random walk (RW);
- RiskMetrics EWMA (J.P. Morgan/Reuters, 1996);
- GARCH(1,1) with Student- t errors, re-estimated monthly;
- the heterogeneous autoregressive (HAR) model of Corsi (2009) with a 7-day week;
- BVIV itself, both raw ($IV_t^2/365$, “IV-raw”) and bias-corrected by regression (“IV”);
- HAR augmented with implied variance (HAR-IV):

$$\ln \overline{RV}_{t+1:t+h} = \beta_0 + \beta_d \ln RV_t + \beta_w \ln \overline{RV}_{t-6:t} + \beta_m \ln \overline{RV}_{t-29:t} + \beta_{iv} \ln (IV_t^2/365) + \varepsilon_{t+h}. \quad (1)$$

HAR and the bias-corrected IV are the special cases without the implied or the realised terms. All regressions are run in logs, as in Andersen et al. (2003), so forecasts $\exp(\hat{y} + s^2/2)$ (with s^2 the training residual variance) are positive by construction and robust to the heavy right tail of realised variance. They are re-estimated daily on an expanding window that contains only targets already observed, with no overlap leakage. We evaluate forecasts with QLIKE, $L(RV, F) = RV/F - \ln(RV/F) - 1$, which is robust to proxy noise (Patton, 2011). We also report the Mincer–Zarnowitz R^2 and Diebold–Mariano tests against HAR with Newey–West variance (Mincer and Zarnowitz, 1969; Diebold and Mariano, 1995).

Results. Table 1 and Figure 1 confirm the pattern documented on Deribit data by Hoang and Baur (2020): implied volatility loses at one day, wins at longer horizons, and combining it with HAR is best at short horizons.

- At one day, HAR-IV is best (median loss ratio 0.95; mean QLIKE 0.219 against 0.231 for HAR) and beats HAR significantly on 26% of paths, losing on 0%. Implied volatility alone is worse (1.18), because a 30-day index cannot track tomorrow’s variance, and GARCH on daily returns lags (1.49), as Bergsli et al. (2022) found for BTC.
- At seven days, HAR-IV remains best (0.89; significant wins on 43% of paths), with the bias-corrected index close behind (0.95).

Model	$h = 1 \text{ day}$			$h = 7 \text{ days}$			$h = 30 \text{ days}$		
	QL ratio	MZ R^2	DM win	QL ratio	MZ R^2	DM win	QL ratio	MZ R^2	DM win
RW	1.25	0.16	0%	1.49	0.17	0%	1.48	0.08	0%
EWMA	1.64	0.08	0%	1.47	0.10	0%	1.59	0.08	0%
GARCH	1.49	0.10	0%	1.26	0.10	0%	1.08	0.04	2%
HAR	1.00	0.20	–	1.00	0.21	–	1.00	0.07	–
IV-raw	1.43	0.18	0%	1.11	0.29	0%	0.96	0.22	6%
IV	1.18	0.19	0%	0.95	0.28	14%	0.85	0.17	26%
HAR-IV	0.95	0.24	26%	0.89	0.30	43%	0.91	0.14	24%

Table 1: Out-of-sample forecast accuracy relative to HAR across 200 test paths. QL ratio: median ratio of a model’s QLIKE to HAR’s (below 1 is better; best in bold). MZ R^2 : mean Mincer–Zarnowitz R^2 . DM win: share of paths on which a Diebold–Mariano test rejects equal accuracy in the model’s favour at the 5% level.




Figure 1: Median QLIKE relative to HAR by model and horizon; the best model at each horizon is highlighted.

- At 30 days, the index is best (0.85), beating HAR significantly on 26% of paths and losing on 0%; HAR-IV follows (0.91).

Because the regressions run in logs, no forecast collapses: across 200 paths and three horizons, HAR-IV’s loss exceeds twice HAR’s in 0 cases. Raw IV is biased upward by the variance risk premium, and the bias correction lowers its 30-day loss ratio from 0.96 to 0.85. For hedge sizing we use the 30-day HAR-IV forecast, because the sizing signal compares a forecast with implied variance, and a forecast built from implied variance alone would add little to that comparison.

4 A BTC book hedged with BVIV perpetuals

P&L. The book is long q BTC at price S_t and long $H_t \geq 0$ BVIV perpetual contracts at mark F_t , each paying \$1 per index point. A position chosen at the close of bar t earns over bar $t+1$:

$$\Pi_{t+1} = q(S_{t+1} - S_t) - \mathbb{1}_{\text{perp}} q S_t f_t \Delta t + H_t(F_{t+1} - F_t) - H_t \psi_t - c|H_{t+1} - H_t| F_{t+1}, \quad (2)$$

where $\Delta t = 1/(365 \cdot 96)$ is the bar length in years, f_t the BTC perpetual’s annualised funding rate, ψ_t the BVIV funding per contract and c the proportional trading cost. A perpetual with funding intensity k trades at $\mathbb{E}_t[\int_0^\infty k e^{-ks} I_{t+s} ds]$, the expected index sampled at a random time (He et al., 2022; Ackerer et al., 2025). In the fast-funding limit the mark equals the index, and funding splits into two parts:

$$\psi_t = \underbrace{\mathbb{E}_t[I_{t+1}] - I_t}_{\text{drift transfer}} + \underbrace{\bar{\kappa} \Delta t}_{\text{carry premium}}. \quad (3)$$

The first term compensates longs exactly for the index’s expected path: after a spike, longs *receive* funding while the index mean-reverts. Only the premium $\bar{\kappa}$ (vol points a year) is a cost, so we report hedge costs as carry premium plus trading costs, not as raw funding.

Hedge ratios. The minimum-variance hedge is $H_t^* = -q S_t \beta_t$, where $\beta_t = \text{Cov}_t(r, \Delta F) / \text{Var}_t(\Delta F)$ is estimated on 15-minute changes with an EWMA (half-life 14 days). We impose $H \geq 0$, since a protective hedge is never short volatility. For crash sizing we also estimate a *downside hedge ratio*

$$\beta_t^- = \mathbb{E}_t[r \Delta F \mid r < 0] / \mathbb{E}_t[\Delta F^2 \mid r < 0]$$

on falling bars only, shrunk towards β_t by its effective sample size, which gives $H_t^- = -qS_t\beta_t^-$. The moments are deliberately not demeaned, so β^- minimises the downside second moment of the hedged book and also offsets the average loss on falling bars.

The VWAP signal. BTC trades 24/7, so we anchor VWAP on a rolling 24-hour window, which avoids an arbitrary session open; Section 6 also tests a UTC-session anchor. Exchange klines give each bar’s exact VWAP as quote volume divided by base volume. For Brownian prices, the gap between the last price and its average over a window w has standard deviation $\sigma\sqrt{w/3}$, so the scale-free breakdown statistic is

$$z_t = \frac{\ln(S_t/\text{VWAP}_t^{24h})}{\hat{\sigma}_t\sqrt{1/3}}, \quad (4)$$

where $\hat{\sigma}_t$ is a causal one-day volatility forecast: an EWMA (half-life one day) of 15-minute squared returns, de-seasonalised by an hour-of-week profile that is estimated on the trailing 60 days, winsorised and re-estimated weekly. Intraday volatility in crypto follows strong and growing periodic patterns (Andersen and Bollerslev, 1997; Hansen et al., 2024; Eross et al., 2019). This intraday estimator, not the daily models of Section 3, sets the band width; the daily forecasts enter the rules only through the ratchet’s size multiplier m_t below. Intraday momentum and reversal both exist in BTC (Shen et al., 2022; Wen et al., 2022), so we treat z as a *state* variable, not a return forecast.

Rules. We compare four families, each with long-only positions, a hedge cap of 50% of BTC notional and a rebalancing band $b = 25\%$. The band follows the no-trade-region logic of Whalley and Wilmott (1997): the hedge is only traded when its target moves more than b away from the current position.

- **Always-on:** $H_t = \lambda H_t^*$, scaled by $\lambda \in \{0.5, 1, 1.5, 2, 3\}$.
- **VWAP-switch:** $H_t = g_t H_t^-$. The gate g_t turns on when $z_t < z_{in}$ and off when $z_t > z_{out}$, after a minimum hold of four hours.
- **VWAP-ratchet:** an overlay o_t jumps to 1 on a breakdown and decays with half-life τ (in days; 96 bars per day):

$$o_t = \begin{cases} 1 & z_t < z_{in} \\ 2^{-1/(96\tau)} o_{t-1} & \text{otherwise,} \end{cases} \quad H_t = \phi H_t^* + o_t \max(H_t^- m_t - \phi H_t^*, 0), \quad (5)$$

where ϕ is a core floor and $m_t = \min(\max(e^{c_t}, 0.5), 1.5)$ an optional size multiplier. The “cheapness” c_t is the log ratio of the previous day’s 30-day HAR-IV variance forecast to implied variance, minus its trailing 180-day median.

- **Oracle:** H_t^- held only during the latent stress regime. It is infeasible and shows what perfect knowledge of the regime is worth; it is not an upper bound for timing, because many crashes start outside stress.

The ratchet’s asymmetry is the point of the design. Re-triggering inside an episode refreshes o_t and rarely forces a trade, so whipsaw around VWAP costs little, whereas the switch pays a round trip every time price crosses the band.

5 Experimental design

Calibrated market. This paper does not analyse live data. We evaluate the rules in a simulated market calibrated to the stylised facts above, which provides a known latent regime, an arbitrage-free perpetual and many independent paths for paired tests. The pipeline ships loaders for live data (Section 7).

- **Variance and returns.** Log-variance is the sum of a slow and a fast Ornstein–Uhlenbeck factor. The fast factor tracks a three-state hidden regime (calm, stress, euphoria; expected durations 30/5/10 days) and receives EGARCH-type shocks (Nelson, 1991), $u_t = \theta_R \eta_t + \gamma(|\eta_t| - \sqrt{2/\pi}) + \sqrt{1 - \theta_R^2 - \gamma^2(1 - 2/\pi)} \varepsilon_t$, where η_t is the standardised return shock. The signed loading $\theta_R = -0.30/-0.60/+0.25$ (calm/stress/euphoria) is negative in calm and stress and positive in euphoria (inverse leverage), and the size loading $\gamma = 0.35$ makes large moves of either sign raise volatility.
- **Jumps and seasonality.** Compound-Poisson jumps cluster in stress and push variance up. Volatility and volume follow hour-of-day and weekend seasonality, and bars are built from one-minute paths so that OHLC and the bar VWAP behave like exchange klines.
- **Index.** bviv is the square root of the model’s own expectation of 30-day average variance, scaled by a variance premium ($\text{IV}^2 \approx e^{0.22} \approx 1.25$ times expected variance), with fixing noise and a persistent sentiment factor (half-life four days). The sentiment shocks load on the same return shock η_t (signed loadings $-0.30/-0.50/+0.20$, size loading 0.20), and this channel carries a large part of the simulated

Moment	Median	5–95%	Empirical anchor
Realised vol, annualised (%)	51.5	[45.1, 59.4]	calibration target ≈ 50
Mean BVIV (vol pts)	57.1	[52.3, 63.1]	35.5 (Aug 2026) to >96 (Feb 2026)
IV – subsequent 30d RV (vol pts)	7.6	[4.0, 10.6]	positive on average
IV ² –RV ² (variance units)	0.06	[0.02, 0.09]	≈ 0.14 in 2017–22 (higher vol)
Share of days IV > RV (%)	78	[66, 87]	VRP positive on average
corr(daily return, Δ IV)	–0.21	[–0.28, –0.12]	sign varies by regime
corr on falling 15m bars	–0.26	[–0.34, –0.20]	IV spikes in sell-offs
corr on rising 15m bars	0.03	[–0.02, 0.09]	IV up in 2023, down in 2025 rallies
Excess kurtosis, daily returns	3.7	[1.7, 10.0]	fat tails
Worst daily loss (%)	15.0	[10.0, 25.3]	fat left tail

Table 2: Simulated moments across 200 test paths of 900 days (15-minute bars).




Figure 2: A representative simulated path: BTC and BVIV on 15-minute bars (4-hour sampling for display), with the latent stress regime shaded. Calendar dates are simulation labels, not history.

spot-vol co-movement. The variance expectation is exact over regime paths: a Feynman–Kac product of matrix exponentials of $Q + \text{diag}(V)$ integrates the fast factor’s regime-driven drift and jump kicks, and the jump variance of returns is mixed with e^{Qs} . A unit test checks it against brute-force simulation.

- *Funding*. The drift transfer in (3) is computed one bar ahead: exactly over regime transitions, by Gauss–Hermite quadrature over jumps, and to second order in the Gaussian shocks, whose one-bar standard deviation is about 0.02. A unit test verifies it against brute-force simulation, so the perpetual offers no free lunch from mean reversion.

Table 2 compares simulated moments with their empirical anchors, and Figure 2 shows a representative path.

Protocol. All signals are causal; unit tests verify that truncating the sample leaves past values of the VWAP signal, the hedge ratios and the regression forecasts unchanged. Rule parameters are selected on 48 *training* paths, separate from the 200 *test* paths reported below. This avoids the data-snooping trap of picking the best of many rules on the evaluation sample (White, 2000; Hansen, 2005; Bailey et al., 2014). The grid contains 87 rules: 72 ratchets ($z_{in} \in \{-1, -1.5, -2\}$, $\tau \in \{0.5, 1, 2, 4\}$ days, $\varphi \in \{0, 0.5, 1\}$, with and without forecast sizing), six switches (z_{in} as above, $z_{out} \in \{0, 0.5\}$), five scaled always-on hedges and four oracle sizes ($\lambda \in \{0.5, 1, 2, 3\}$). Only the switches and ratchets are candidates for selection. The selection criterion is the median reduction in expected shortfall ($ES_{97.5}$) of daily returns, measured *net* of carry and trading costs. Hedging is evaluated from day 270, once every forecast and signal exists, and forecasts from day 365. Costs are 6 bp fee plus 10 bp slippage per side on traded notional, and the carry is 6 vol points a year.

6 Results

Selected rules. On the training paths, the best ratchet keeps the full minimum-variance hedge as its core ($\varphi = 1$), triggers at $z < -1.5$ and decays with a half-life of 4 days. It runs with forecast-based sizing, but this barely matters: on the test paths its unsized twin differs by +0.02 percentage points (pp) of ES reduction

Rule	ES red.	Var red.	Max DD	Tail offset	Carry	Trading	Turnover	Time on
Unhedged	0.0	0.0	49.5	0.0	0.00	0.00	0.0	0
Always-on (MV)	4.2	4.1	48.7	5.9	0.74	0.14	0.9	93
Always-on, 1.5×	4.9	3.1	49.5	8.8	1.11	0.20	1.3	93
VWAP-switch	2.8	1.9	50.9	4.9	0.38	3.94	24.6	35
VWAP-ratchet	4.9	3.9	48.9	7.8	0.96	0.23	1.4	98
Ratchet, placebo timing	4.9	3.8	48.8	7.7	0.96	0.23	1.5	98
Oracle (infeasible)	2.3	2.3	48.1	3.5	0.11	0.27	1.7	9

Table 3: Hedging a long-BTC spot book on 200 test paths (medians; % unless stated). ES/Var red.: reduction in daily ES_{97.5} and variance relative to unhedged. Max DD: maximum drawdown. Tail offset: hedge P&L net of costs as a share of BTC losses on the worst 5% of days. Carry and Trading: costs in % of BTC notional a year. Turnover: bviv-perpetual notional traded per year as a multiple of BTC notional. Time on: share of bars with a hedge. Placebo timing: the selected ratchet with its trigger series circularly shifted (mean of 5 shifts).




Figure 3: The deepest fresh VWAP breakdown on the representative test path (a trigger after at least three trigger-free days). Top: 15-minute close, rolling 24-hour VWAP and the breakdown band (orange dots: trigger bars). Bottom: hedge sizes. The switch buys protection only once the breakdown is under way and sells it as price recovers, paying a round trip. The selected ratchet barely moves: with a 4-day half-life and the 25% rebalancing band, earlier triggers keep it near its upper level, so in effect it is a larger always-on hedge, which is the size effect identified by the placebo test.

(95% CI $[-0.09, +0.13]$). The best switch enters at $z < -1$ and exits at $z > 0.5$.

Hedging performance. Table 3 summarises the test results, and Figure 3 shows the rules on one episode.

- *Always-on.* The minimum-variance hedge reduces ES by 4.2% and variance by 4.1% with a notional of only 7.0% of the BTC position. A volatility index is a partial hedge for a long-BTC book: the simulated correlation between daily returns and bviv changes is -0.21 (on 15-minute bars, -0.26 when BTC falls and 0.03 when it rises), and a static minimum-variance hedge removes only the squared correlation's share of the variance (Ederington, 1979).
- *VWAP-switch.* The switch is the instructive failure. It trades 24.6 times the BTC notional per year, and its ES reduction trails the always-on hedge by 1.67 pp (paired 95% CI $[-2.00, -1.36]$). 86% of its openings are not followed by one of the worst 5% of days within 24 hours, and each costs a round trip.
- *VWAP-ratchet.* The ratchet reduces ES by 4.9% at 1.19% a year (always-on: 0.88%). Paired across test paths, it beats the always-on hedge by +0.53 pp (95% CI $[+0.33, +0.72]$; higher on 66% of paths) and offsets +1.9 pp more of the tail losses, while its variance reduction is 0.32 pp lower.
- *Oracle.* Holding the hedge only in the latent stress regime reduces ES by just 2.3%: 62% of the unhedged book's worst 2.5% of days contain no stress-regime bar (first 24 test paths), so regime timing without a core hedge leaves the book exposed.




Figure 4: Tail protection versus cost on the test paths (medians). Blue: the always-on hedge scaled from $0.5\times$ to $3\times$. Dots: every candidate VWAP rule; diamonds: rules selected on the training paths, with interquartile ranges of the ES reduction.

Timing or size? A placebo test. The ratchet holds more hedge on average than the always-on rule (9.1% versus 7.0% of notional), so its gain could come from size alone. To separate the two, we circularly shift each path’s trigger series by 35 to 259 days, in multiples of a week. This keeps the number, clustering and hour-of-week profile of the triggers but breaks their link to the market. For the selected ratchet the placebo is just as good: it reduces ES by 4.9%, and the ratchet beats it by only +0.05 pp (95% CI $[-0.03, +0.13]$). Scaling the static hedge by 1.5 achieves the same ES reduction (4.9% at 1.32%), although with less variance reduction (3.1% versus 3.9%).

The grid tells a more nuanced story. We define a rule’s timing value as its ES reduction minus the mean over its placebos, and compute it for all 72 ratchet configurations on the first 64 test paths with 3 shifts each. It averages +0.31 pp and ranges from -0.15 to $+1.36$ pp. Of the configurations, 31 beat their placebo at the 5% level and none loses to it significantly. VWAP breakdowns therefore do carry information about tail days, but how much depends on the core hedge: without a core ($\varphi = 0$) timing adds +0.71 pp on average (18 of 24 configurations significant), with half a core +0.22 pp, and with the full core +0.01 pp. Rules that rely on timing protect less (median ES reduction 3.6% without a core, against 4.9% with one), and the correlation between a configuration’s timing value and its ES reduction is -0.40 . The most timing-dependent rule reaches 4.1% at a cost of 4.2% a year, against 4.6% for the always-on hedge on the same paths. Being hedged when the breakdown comes matters more than reacting to it.

Sizing. Figure 4 places every rule in cost-protection space. Along the scaled static hedge, ES reduction peaks near $1.5\times$ (4.9%) and then falls (4.5% at $2\times$, 1.8% at $3\times$). The extra contracts add volatility-of-volatility risk faster than they add crash convexity. Only 9 of the 72 ratchet configurations (12%; 7 of them with a full core) beat every static hedge that costs no more. The best do so through a better-shaped size profile, not timing: they buy the ES reduction of a $1.5\times$ static hedge with more variance reduction and at lower cost. The economically first-order choice is how much BVIV to hold.

Robustness. Table 4 varies one assumption at a time.

- *Costs and carry.* Doubling or halving trading costs ($\times 2 / \times 0.5$), or setting the carry to zero or twelve vol points, leaves the ratchet’s edge over the always-on hedge between +0.63 and +0.67 pp, significant throughout.
- *Spot-vol regime.* The regime matters more than any rule. Under VIX-like dynamics ($\theta_R = -0.50 / -0.75 / -0.10$, sentiment $-0.45 / -0.60 / -0.10$) the always-on hedge cuts ES by 9.7% and the ratchet’s edge is +0.74 pp. Under inverse leverage ($\theta_R = 0 / -0.40 / +0.50$, sentiment $0 / -0.30 / +0.40$) the hedge is nearly useless (0.5%), and the edge shrinks to +0.18 pp (CI $[-0.02, +0.39]$).
- *Crash clustering.* When crashes cluster in the stress regime, the edge roughly halves to +0.34 pp (CI $[0.00, +0.68]$), and the ratchet’s median ES reduction falls below the always-on hedge’s (5.1% versus 5.3%). The larger post-breakdown position then also exposes the book to volatility crushes. On the hedged book’s own worst 2.5% of days (first 24 test paths), the ratchet’s hedge leg returns -0.04% of BTC notional,

Scenario	ES red. (%)		Ratchet – Always		Switch ES red.	Cost (% p.a.)	
	Always	Ratchet	mean [95% CI]	> 0		Always	Ratchet
Baseline	4.3	5.1	+0.66 [+0.39, +0.94]	71%	3.6	0.87	1.22
Trading costs $\times 2$	4.3	5.1	+0.63 [+0.37, +0.91]	71%	3.4	1.01	1.46
Trading costs $\times 0.5$	4.3	5.2	+0.67 [+0.41, +0.95]	71%	3.7	0.80	1.10
Carry 0 vol pts	4.3	5.2	+0.67 [+0.40, +0.95]	71%	3.7	0.13	0.24
Carry 12 vol pts	4.3	5.1	+0.65 [+0.38, +0.93]	71%	3.6	1.61	2.20
VIX-like spot-vol	9.7	10.2	+0.74 [+0.47, +1.03]	69%	6.3	1.33	1.64
Inverse leverage	0.5	0.7	+0.18 [−0.02, +0.39]	58%	0.5	0.25	0.57
Crashes cluster in stress	5.3	5.1	+0.34 [−0.00, +0.68]	61%	3.1	0.85	1.18
Session-anchored VWAP	4.3	5.4	+0.73 [+0.44, +1.02]	73%	3.4	0.87	1.25
Perpetual BTC book	4.3	5.1	+0.64 [+0.38, +0.92]	71%	3.6	0.87	1.22

Table 4: Robustness on 96 fresh paths per scenario (a separate seed block, common across scenarios), with the selected rules held fixed. ES red.: median reduction in daily $ES_{97.5}$. Ratchet – Always: mean paired difference in percentage points, with a 95% bootstrap CI and the share of paths with a positive difference. Cost: carry premium plus trading costs.

against +0.16% for the always-on hedge (baseline: +0.13% and +0.28%). As with equity puts, badly timed protection can deepen losses (Israelov, 2019).

- *VWAP anchor and book.* A UTC-session anchor gives +0.73 pp. Hedging a perpetual instead of spot gives +0.64 pp: BTC funding shifts the book’s mean return but not its tail co-movement with bv_{IV} .

7 Implementation on live data

The pipeline (`scripts/run_paper.py --live`) maps live data into the simulator’s schema, so forecasting, signals and hedging share one code path. It backtests the selected rules and runs the timing placebo on the live path, comparing the ratchet with every weekly shift of its triggers. Its sources are BTCUSDT 15-minute klines from the Binance bulk archive (quote volume over volume gives each bar’s exact VWAP) and BTC perpetual funding; index, mark and funding of the Bitfinex bv_{IV} perpetual since April 2024; and bv_{IV} bars from the Volmex API, or Deribit DVOL as a whole-sample proxy (the series are not spliced). Without perpetual data, funding is set to the constant carry, so hedge P&L then includes the index’s predictable mean reversion, which a tradable perpetual would price out. The bv_{IV} perpetuals are young and thin, so the costs in Table 4 should be stressed further for large books. Oracle risk is real on builder-deployed markets: in July 2026 a single-share pre-market print, about 30% below the prior close, moved a HIP-3 equity perpetual on Hyperliquid by about 19% and triggered large liquidations. Bitfinex’s funding process can briefly pause trading, so a hedge should not depend on one venue in a crash. Finally, the 15-minute boundary attracts algorithmic order flow (Kim and Hansen, 2026), so hedge orders should be worked inside the bar, for example by VWAP slicing (Berkowitz et al., 1988; Konishi, 2002; Białkowski et al., 2008), not sent at the close (Madhavan, 2002).

Limitations. The evidence is simulation-based. The simulator links VWAP breakdowns to crashes only through regime drift and volatility clustering. Real liquidation cascades may make intraday breakdowns more informative, for instance through intraday momentum (Shen et al., 2022), or less. Our finding that a core hedge makes timing redundant is therefore a statement about this calibrated market, not a theorem. On live data, the placebo test of Section 6 is the decisive check; the live mode runs it on the single available path, judging the ratchet against the distribution of all weekly shifts of its triggers. The forecast ranking, too, needs live confirmation, since the simulated index is a near-rational forecast by construction, although it already matches the literature (Hoang and Baur, 2020; Bergsli et al., 2022).

8 Conclusion

bv_{IV} perpetuals make BTC implied volatility a hedging instrument, but a partial one. Four lessons follow for this calibrated market, which the live pipeline can test. (1) Forecast BTC volatility with HAR augmented with implied variance at horizons up to a week, and with the bias-corrected implied index at a month; run the regressions in logs. (2) Size a bv_{IV} hedge with the minimum-variance ratio (up to about 1.5 times it when targeting expected shortfall) and keep it on. (3) Do not switch the hedge around VWAP, because turnover eats the protection. (4) Treat a 15-minute VWAP breakdown as real but secondary information: it beats random timing only without a core hedge; with one, a ratchet’s gain is size, not timing. Any intraday timing rule should beat a placebo with the same trigger statistics before it earns its complexity.

References

- Damien Ackerer, Julien Hugonnier, and Urban Jermann. Perpetual futures pricing. *Mathematical Finance*, 2025. doi: 10.1111/mafi.70018.
- Carol Alexander and Arben Imeraj. The Bitcoin VIX and its variance risk premium. *Journal of Alternative Investments*, 23(4):84–109, 2021. doi: 10.3905/jai.2020.1.112.
- Carol Alexander, Dimitris Korovilas, and Julia Kapraun. Diversification with volatility products. *Journal of International Money and Finance*, 65: 213–235, 2016. doi: 10.1016/j.jimonfin.2015.10.005.
- Caio Almeida, Maria Grith, Ratmir Miftachov, and Zijin Wang. Risk premia in the Bitcoin market. arXiv preprint 2410.15195, revised August 2025, 2024.
- Torben G. Andersen and Tim Bollerslev. Intraday periodicity and volatility persistence in financial markets. *Journal of Empirical Finance*, 4(2–3): 115–158, 1997. doi: 10.1016/S0927-5398(97)00004-2.
- Torben G. Andersen, Tim Bollerslev, Francis X. Diebold, and Paul Labys. Modeling and forecasting realized volatility. *Econometrica*, 71(2):579–625, 2003. doi: 10.1111/1468-0262.00418.
- David H. Bailey, Jonathan M. Borwein, Marcos López de Prado, and Qiji Jim Zhu. Pseudo-mathematics and financial charlatanism: The effects of backtest overfitting on out-of-sample performance. *Notices of the American Mathematical Society*, 61(5):458–471, 2014. doi: 10.1090/noti1105.
- Dirk G. Baur and Thomas Dimpfl. Asymmetric volatility in cryptocurrencies. *Economics Letters*, 173:148–151, 2018. doi: 10.1016/j.econlet.2018.10.008.
- Lykke Ø. Bergsli, Andrea F. Lind, Peter Molnár, and Michał Polasik. Forecasting volatility of Bitcoin. *Research in International Business and Finance*, 59:101540, 2022. doi: 10.1016/j.ribaf.2021.101540.
- Stephen A. Berkowitz, Dennis E. Logue, and Eugene A. Noser, Jr. The total cost of transactions on the NYSE. *Journal of Finance*, 43(1):97–112, 1988. doi: 10.1111/j.1540-6261.1988.tb02591.x.
- Jędrzej Białkowski, Serge Darolles, and Gaëlle Le Fol. Improving VWAP strategies: A dynamic volume approach. *Journal of Banking & Finance*, 32(9):1709–1722, 2008. doi: 10.1016/j.jbankfin.2007.09.023.
- Elie Bouri, Georges Azzi, and Anne Haubo Dyhrberg. On the return-volatility relationship in the Bitcoin market around the price crash of 2013. *Economics: The Open-Access, Open-Assessment E-Journal*, 11(2017-2), 2017. doi: 10.5018/economics-ejournal.ja.2017-2.
- Peter Carr and Liuren Wu. Variance risk premiums. *Review of Financial Studies*, 22(3):1311–1341, 2009. doi: 10.1093/rfs/hhn038.
- CME Group. CME Group to launch Bitcoin volatility futures contracts. Press release, May, 2026.
- CoinDesk. Bitcoin volatility index’s correlation with price flips positive, boosts appeal of bullish call options. CoinDesk, 4 April, 2023.
- CoinDesk. Bitfinex derivatives users can now place bets on bitcoin and ether implied volatility. CoinDesk, 3 April, 2024.
- CoinDesk. Volmex’s bitcoin and ether volatility futures top \$10M volume since debut as traders look beyond price. CoinDesk, 25 July, 2025a.
- CoinDesk. Bitcoin’s record rally from \$70K to \$118K: A tale of transition from wild west to wall street-like dynamics. CoinDesk, 11 July, 2025b.
- CoinDesk. Perp futures linked to ‘bitcoin VIX’ debut on Hyperliquid. CoinDesk, 21 September, 2026.
- Fulvio Corsi. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2):174–196, 2009. doi: 10.1093/jfinec/nbp001.
- Kresimir Demeterfi, Emanuel Derman, Michael Kamal, and Joseph Zou. A guide to volatility and variance swaps. *Journal of Derivatives*, 6(4):9–32, 1999. doi: 10.3905/jod.1999.319129.
- Deribit. DVOL – Deribit implied volatility index. Deribit Insights, exchange documentation, 2021.
- Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995. doi: 10.1080/07350015.1995.10524599.
- Louis H. Ederington. The hedging performance of the new futures markets. *Journal of Finance*, 34(1):157–170, 1979. doi: 10.1111/j.1540-6261.1979.tb02077.x.
- Robert F. Engle. Dynamic conditional correlation: A simple class of multivariate generalized autoregressive conditional heteroskedasticity models. *Journal of Business & Economic Statistics*, 20(3):339–350, 2002. doi: 10.1198/073500102288618487.
- Andrea Eross, Frank McGroarty, Andrew Urquhart, and Simon Wolfe. The intraday dynamics of Bitcoin. *Research in International Business and Finance*, 49:71–81, 2019. doi: 10.1016/j.ribaf.2019.01.008.
- Peter Reinhard Hansen. A test for superior predictive ability. *Journal of Business & Economic Statistics*, 23(4):365–380, 2005. doi: 10.1198/073500105000000063.
- Peter Reinhard Hansen, Chan Kim, and Wade Kimbrough. Periodicity in cryptocurrency volatility and liquidity. *Journal of Financial Econometrics*, 22:224–251, 2024. doi: 10.1093/jfinec/nbac034.
- Songrun He, Asaf Manela, Omri Ross, and Victor von Wachter. Fundamentals of perpetual futures. arXiv preprint 2212.06888, 2022.
- Lai T. Hoang and Dirk G. Baur. Forecasting Bitcoin volatility: Evidence from the options market. *Journal of Futures Markets*, 40(10):1584–1602, 2020. doi: 10.1002/fut.22144.
- Jing-Zhi Huang, Jun Ni, and Li Xu. Leverage effect in cryptocurrency markets. *Pacific-Basin Finance Journal*, 73:101773, 2022. doi: 10.1016/j.pacfin.2022.101773.
- Hyperliquid. Funding; HIP-3 builder-deployed perpetuals. Hyperliquid documentation, 2026.
- Roni Israelov. Pathetic protection: The elusive benefits of protective puts. *Journal of Alternative Investments*, 21(3):6–33, 2019. doi: 10.3905/jai.2018.1.066.
- Leland L. Johnson. The theory of hedging and speculation in commodity futures. *Review of Economic Studies*, 27(3):139–151, 1960. doi: 10.2307/2296076.
- J.P. Morgan/Reuters. *RiskMetrics – Technical Document*. J.P. Morgan/Reuters, New York, 4th edition, 1996.
- Chan Kim and Peter Reinhard Hansen. The quarter-hour effect: Periodic algorithmic trading and return predictability in cryptocurrency futures. arXiv preprint 2607.09426, 2026.
- Hizuru Konishi. Optimal slice of a VWAP trade. *Journal of Financial Markets*, 5(2):197–221, 2002. doi: 10.1016/S1386-4181(01)00023-4.
- Kenneth F. Kroner and Jahangir Sultan. Time-varying distributions and dynamic hedging with foreign currency futures. *Journal of Financial and Quantitative Analysis*, 28(4):535–551, 1993. doi: 10.2307/2331164.
- Lily Y. Liu, Andrew J. Patton, and Kevin Sheppard. Does anything beat 5-minute RV? A comparison of realized measures across multiple asset classes. *Journal of Econometrics*, 187(1):293–311, 2015. doi: 10.1016/j.jeconom.2015.02.008.
- Ananth Madhavan. VWAP strategies. *Trading*, 2002(1):32–39, 2002.
- Jacob A. Mincer and Victor Zarnowitz. The evaluation of economic forecasts. In Jacob A. Mincer, editor, *Economic Forecasts and Expectations: Analysis of Forecasting Behavior and Performance*, pages 3–46. NBER, 1969.
- Daniel B. Nelson. Conditional heteroskedasticity in asset returns: A new approach. *Econometrica*, 59(2):347–370, 1991. doi: 10.2307/2938260.
- Andrew J. Patton. Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1):246–256, 2011. doi: 10.1016/j.jeconom.2010.03.034.
- Dehua Shen, Andrew Urquhart, and Pengfei Wang. Bitcoin intraday time series momentum. *Financial Review*, 57(2):319–344, 2022. doi: 10.1111/fire.12290.
- Edward Szado. VIX futures and options: A case study of portfolio diversification during the 2008 financial crisis. *Journal of Alternative Investments*, 12(2):68–85, 2009. doi: 10.3905/JAI.2009.12.2.068.
- Tetsuya Takaishi. Time-varying properties of asymmetric volatility and multifractality in Bitcoin. *PLoS ONE*, 16(2):e0246209, 2021. doi: 10.1371/journal.pone.0246209.
- Volmex Labs. Volmex implied volatility indices: Methodology. White paper, <https://volmex.finance/Volmex-IV-paper.pdf>, n.d. Accessed October 2026.
- Zhuzhu Wen, Elie Bouri, Yahua Xu, and Yang Zhao. Intraday return predictability in the cryptocurrency markets: Momentum, reversal, or both. *North American Journal of Economics and Finance*, 62:101733, 2022. doi: 10.1016/j.najef.2022.101733.
- Robert E. Whaley. Trading volatility: At what cost? *Journal of Portfolio Management*, 40(1):95–108, 2013. doi: 10.3905/jpm.2013.40.1.095.
- A. Elizabeth Whalley and Paul Wilmott. An asymptotic analysis of an optimal hedging model for option pricing with transaction costs. *Mathematical Finance*, 7(3):307–324, 1997. doi: 10.1111/1467-9965.00034.
- Halbert White. A reality check for data snooping. *Econometrica*, 68(5): 1097–1126, 2000. doi: 10.1111/1468-0262.00152.
- Holbrook Working. Note on the correlation of first differences of averages in a random chain. *Econometrica*, 28(4):916–918, 1960. doi: 10.2307/1907574.